Decision (date)	Rationale
Maintenance-only frontend (16 Oct)	The product is maintenance diagnostics; the legacy software screens stay in the repository but are unreachable, so they cannot be mistaken for features
Frontend reflects backend outputs (16 Oct)	The interface may add labels and warnings but must never hide, alter or invent a value the backend returned
Review actions are server-confirmed (16 Oct)	Accept and reject change real data, so the screen changes only after the server confirms; no keyboard shortcut can trigger them
Centralised staging and CI-gated release (16 Oct)	One shared staging API for the whole team; only a commit that passed full CI on `main` is deployed, with a recorded rollback
Honest analysis state (17 Oct)	Analysis state is stored separately from confidence; "unscored" is never shown as a confident result
Fleet-wide paginated records (17 Oct)	Replaced a 119-request fan-out with one paginated, filterable endpoint
Opt-in cost cascade (17 Oct)	Pre-registered, off by default: a deterministic text screen before the model, and adaptive voting, to cut requests without tuning against the evaluation
Diagnostic copilot scope waiver (16 Oct)	A named, limited exception to file ownership so the copilot could be built across route, retrieval and UI without touching the conflict engine, ingest or CI
Graph shows one aircraft (17 Oct)	The whole graph is too dense to read; one aircraft's stored subgraph is readable and every drawn edge is a stored edge
Knowledge bank and agent protocol (16 Oct)	A written, versioned memory of context, flows and decisions that every contributor and every AI assistant must read first and update in the same change