RQ	Question	Planned evidence
RQ1	Does graph-constrained retrieval beat keyword retrieval at finding relevant cases?	Recall at k and mean reciprocal rank on the 141 FAA pairs, same queries for both methods
RQ2	Does hybrid retrieval improve ranking over semantic alone?	Same metrics, with an ablation of each signal
RQ3	Does correction-aware retrieval reduce misleading recommendations?	Using the FAA pairs in which the follow-up names a different component (92 of 141): does the system surface the later component for the earlier symptom?
RQ4	Can abstention reduce unsupported suggestions?	Coverage against error rate as the evidence threshold changes
RQ5	Does case-based training improve learners' results?	Planned with the training module in FYP-2